

# TABLE OF CONTENTS

- ◆ INTRODUCTION
- ◆ INSERTION
- ◆ SEARCHING
- ◆ DELETION
- ◆ FIND min AND max
- ◆ TREE TRAVERSALS
  - ◇ Pre order
  - ◇ Post order
  - ◇ In order
- ◆ BREADTH FIRST

## INTRODUCTION

Binary search trees (BSTs) are very simple to understand. We start with a root node with value  $x$ , where the left subtree of  $x$  contains nodes with values  $< x$  and the right subtree contains nodes whose values are  $\geq x$ . Each node follows the same rules with respect to nodes in their left and right subtrees.

BSTs are of interest because they have operations which are favourably fast: insertion, look up, and deletion can all be done in  $O(\log n)$  time. It is important to note that the  $O(\log n)$  times for these operations can only be attained if the BST is reasonably balanced; for a tree data structure with self balancing properties see AVL tree defined in §7).

## INSERTION

```

1) algorithm Insert(value)
2)   Pre: value has passed custom type checks for type T
3)   Post: value has been placed in the correct location in the tree
4)   if root =  $\emptyset$ 
5)     root  $\leftarrow$  node(value)
6)   else
7)     InsertNode(root, value)
8)   end if
9) end Insert
  
```

```

1) algorithm InsertNode(current, value)
2)   Pre: current is the node to start from
3)   Post: value has been placed in the correct location in the tree
4)   if value < current.Value
5)     if current.Left =  $\emptyset$ 
6)       current.Left  $\leftarrow$  node(value)
7)     else
8)       InsertNode(current.Left, value)
9)     end if
10)  else
11)    if current.Right =  $\emptyset$ 
12)      current.Right  $\leftarrow$  node(value)
13)    else
14)      InsertNode(current.Right, value)
15)    end if
16)  end if
17) end InsertNode
  
```

The insertion algorithm is split for a good reason. The first algorithm (non-recursive) checks a very core base case - whether or not the tree is empty. If the tree is empty then we simply create our root node and finish. In all other cases we invoke the recursive *InsertNode* algorithm which simply guides us to the first appropriate place in the tree to put *value*. Note that at each stage we perform a binary chop: we either choose to recurse into the left subtree or the right by comparing the new value with that of the current node. For any totally ordered type, no value can simultaneously satisfy the conditions to place it in both subtrees.

## Searching

We have talked previously about insertion, we go either left or right with the right subtree containing values that are  $\geq x$  where  $x$  is the value of the node we are inserting. When searching the rules are made a little more atomic and at any one time we have four cases to consider:

1. the  $root = \emptyset$  in which case *value* is not in the BST; or
2.  $root.Value = value$  in which case *value* is in the BST; or
3.  $value < root.Value$ , we must inspect the left subtree of *root* for *value*; or
4.  $value > root.Value$ , we must inspect the right subtree of *root* for *value*.

```

1) algorithm Contains(root, value)
2)   Pre: root is the root node of the tree, value is what we would like to locate
3)   Post: value is either located or not
4)   if  $root = \emptyset$ 
5)     return false
6)   end if
7)   if  $root.Value = value$ 
8)     return true
9)   else if  $value < root.Value$ 
10)    return Contains(root.Left, value)
11)  else
12)    return Contains(root.Right, value)
13)  end if
14) end Contains

```

## Deletion

Removing a node from a BST is fairly straightforward, with four cases to consider:

1. the value to remove is a leaf node; or
2. the value to remove has a right subtree, but no left subtree; or
3. the value to remove has a left subtree, but no right subtree; or
4. the value to remove has both a left and right subtree in which case we promote the largest value in the left subtree.

There is also an implicit fifth case whereby the node to be removed is the only node in the tree. This case is already covered by the first, but should be noted as a possibility nonetheless.

Of course in a BST a value may occur more than once. In such a case the first occurrence of that value in the BST will be removed.




Figure 3.2: binary search tree deletion cases

```

1) algorithm FindNode(root, value)
2)   Pre: value is the value of the node we want to find the parent of
3)   root is the root node of the BST
4)   Post: a reference to the node of value if found; otherwise ∅
5)   if root = ∅
6)     return ∅
7)   end if
8)   if root.Value = value
9)     return root
10)  else if value < root.Value
11)    return FindNode(root.Left, value)
12)  else
13)    return FindNode(root.Right, value)
14)  end if
15) end FindNode
  
```

```

1) algorithm FindParent(value, root)
2)   Pre: value is the value of the node we want to find the parent of
3)   root is the root node of the BST and is ! = ∅
4)   Post: a reference to the parent node of value if found; otherwise ∅
5)   if value = root.Value
6)     return ∅
7)   end if
8)   if value < root.Value
9)     if root.Left = ∅
10)      return ∅
11)    else if root.Left.Value = value
12)      return root
13)    else
14)      return FindParent(value, root.Left)
15)    end if
16)  else
17)    if root.Right = ∅
18)      return ∅
19)    else if root.Right.Value = value
20)      return root
21)    else
22)      return FindParent(value, root.Right)
23)    end if
24)  end if
25) end FindParent
  
```

Missing root parameter in remove function, because we will need it in FindNode

```

1) algorithm Remove(value)
2)   Pre: value is the value of the node to remove, root is the root node of the BST
3)   Count is the number of items in the BST
4)   Post: node with value is removed if found in which case yields true, otherwise false
5)   nodeToRemove  $\leftarrow$  FindNode(value)
6)   if nodeToRemove =  $\emptyset$ 
7)     return false // value not in BST
8)   end if
9)   parent  $\leftarrow$  FindParent(value)
10)  if Count = 1
11)    root  $\leftarrow \emptyset$  // we are removing the only node in the BST
12)  else if nodeToRemove.Left =  $\emptyset$  and nodeToRemove.Right = null
13)    // case #1
14)    if nodeToRemove.Value < parent.Value
15)      parent.Left  $\leftarrow \emptyset$ 
16)    else
17)      parent.Right  $\leftarrow \emptyset$ 
18)    end if
19)  else if nodeToRemove.Left =  $\emptyset$  and nodeToRemove.Right  $\neq \emptyset$ 
20)    // case #2
21)    if nodeToRemove.Value < parent.Value
22)      parent.Left  $\leftarrow$  nodeToRemove.Right
23)    else
24)      parent.Right  $\leftarrow$  nodeToRemove.Right
25)    end if
26)  else if nodeToRemove.Left  $\neq \emptyset$  and nodeToRemove.Right =  $\emptyset$ 
27)    // case #3
28)    if nodeToRemove.Value < parent.Value
29)      parent.Left  $\leftarrow$  nodeToRemove.Left
30)    else
31)      parent.Right  $\leftarrow$  nodeToRemove.Left
32)    end if
33)  else
34)    // case #4
35)    largestValue  $\leftarrow$  nodeToRemove.Left
36)    while largestValue.Right  $\neq \emptyset$ 
37)      // find the largest value in the left subtree of nodeToRemove
38)      largestValue  $\leftarrow$  largestValue.Right
39)    end while
40)    // set the parents' Right pointer of largestValue to  $\emptyset$ 
41)    FindParent(largestValue.Value).Right  $\leftarrow \emptyset$ 
42)    nodeToRemove.Value  $\leftarrow$  largestValue.Value
43)  end if
44)  Count  $\leftarrow$  Count - 1
45) return true
end Remove

```



Astute readers will have noticed that the *FindNode* algorithm is exactly the same as the *Contains* algorithm (defined in §3.2) with the modification that we are returning a reference to a node not *true* or *false*. Given *FindNode*, the easiest way of implementing *Contains* is to call *FindNode* and compare the return value with  $\emptyset$ .

## Finding the smallest and largest values

```

1) algorithm FindMin(root)
2)   Pre: root is the root node of the BST
3)   root  $\neq \emptyset$ 
4)   Post: the smallest value in the BST is located
5)   if root.Left =  $\emptyset$ 
6)     return root.Value
7)   end if
8)   FindMin(root.Left)
9) end FindMin

```

```

1) algorithm FindMax(root)
2)   Pre: root is the root node of the BST
3)   root  $\neq \emptyset$ 
4)   Post: the largest value in the BST is located
5)   if root.Right =  $\emptyset$ 
6)     return root.Value
7)   end if
8)   FindMax(root.Right)
9) end FindMax

```

# Tree Traversals

## Preorder

Pre-order Traversal is done by **visiting the root node first**, then **recursively do a pre-order traversal of the left subtree**, followed by a **recursive pre-order traversal of the right subtree**. It's **used for creating a copy of the tree**, prefix notation of an expression tree, etc.

When using the preorder algorithm, **you visit the root first, then traverse the left subtree and finally traverse the right subtree**. An example of preorder traversal is shown in Figure 3.3.




Figure 3.3: Preorder visit binary search tree example

```
1) algorithm Preorder(root)
2)   Pre: root is the root node of the BST
3)   Post: the nodes in the BST have been visited in preorder
4)   if root  $\neq \emptyset$ 
5)     yield root.Value
6)     Preorder(root.Left)
7)     Preorder(root.Right)
8)   end if
9) end Preorder
```

## Postorder

the **value of the node is yielded after** traversing both subtrees.






Figure 3.4: Postorder visit binary search tree example


```

1) algorithm Postorder(root)
2)   Pre: root is the root node of the BST
3)   Post: the nodes in the BST have been visited in postorder
4)   if root  $\neq \emptyset$ 
5)     Postorder(root.Left)
6)     Postorder(root.Right)
7)     yield root.Value
8)   end if
9) end Postorder

```

## Inorder

the value of the current node is yielded in between traversing the left subtree and the right subtree.



```

1) algorithm Inorder(root)
2)   Pre: root is the root node of the BST
3)   Post: the nodes in the BST have been visited in inorder
4)   if root  $\neq \emptyset$ 
5)     Inorder(root.Left)
6)     yield root.Value
7)     Inorder(root.Right)
8)   end if
9) end Inorder


```

One of the beauties of inorder traversal is that values are yielded in their comparison order. In other words, when traversing a populated BST with the inorder strategy, the yielded sequence would have property  $x_i \leq x_{i+1} \forall i$ .

## Breadth First

Traversing a tree in breadth first order **yields the values of all nodes of a particular depth in the tree before any deeper ones**. In other words, given a depth **d** we would visit the values of all nodes at **d** in a **left to right** fashion, then we would proceed to **d + 1** and so on until we had no more nodes to visit.

**Traditionally** breadth first traversal is implemented using **a list** (vector, resizable array, etc) to **store the values of the nodes visited** in breadth first order and then a **queue** to store those **nodes that have yet to be visited**



```

1) algorithm BreadthFirst(root)
2)   Pre: root is the root node of the BST
3)   Post: the nodes in the BST have been visited in breadth first order
4)   q  $\leftarrow$  queue
5)   while root  $\neq \emptyset$ 
6)     yield root.Value
7)     if root.Left  $\neq \emptyset$ 
8)       q.Enqueue(root.Left)
9)     end if
10)    if root.Right  $\neq \emptyset$ 
11)      q.Enqueue(root.Right)
12)    end if
13)    if !q.IsEmpty()
14)      root  $\leftarrow$  q.Dequeue()
15)    else
16)      root  $\leftarrow \emptyset$ 
17)    end if
18)  end while
19) end BreadthFirst

```

## Summary

A binary search tree is a good solution when you need to represent types that are ordered according to some custom rules inherent to that type. With logarithmic insertion, lookup, and deletion it is very efficient. Traversal remains linear, but there are many ways in which you can visit the nodes of a tree. Trees are recursive data structures, so typically you will find that many algorithms that operate on a tree are recursive.

The run times presented in this chapter are based on a pretty big assumption - that the binary search tree's left and right subtrees are reasonably balanced. We can only attain logarithmic run times for the algorithms presented earlier when this is true. A binary search tree does not enforce such a property, and the run times for these operations on a pathologically unbalanced tree become linear: such a tree is effectively just a linked list. Later in §7 we will examine an AVL tree that enforces self-balancing properties to help attain logarithmic run times.